Your grade: 100%

Your latest: **100%** • Your highest: **100%** • To pass you need at least 60%. We keep your highest score.

Next item

1. Cassandra is the best fit for use cases that have which characteristic? 1 / 1 point

- (●) Always available
- ( ) Queries and joins
- ( ) Frequent schema changes
- ( ) Focused on search

Correct
Correct! Cassandra is the best fit for globally “always available” types of online applications. Cassandra is also a good choice for strong transactions, and time-series use cases.

2. Which of the following statements best describes the distributed data architecture of Cassandra? 1 / 1 point

- ( ) Data is replicated across all nodes
- ( ) Data is distributed in a centralized data warehouse
- ( ) Data is distributed for load balancing
- (●) Data is evenly distributed across nodes

Correct
Correct! Cassandra’s architecture distributes data across nodes.

3. Which phrase describes a consequence of Cassandra’s availability? 1 / 1 point

- (●) Possible inconsistency of the changed data
- ( ) Irresolvable data conflicts
- ( ) Strong consistency
- ( ) Little to zero consistency

Correct
Correct! Even if you lose a part of your cluster, nodes remain available to answer the service request; however, the returned data might be inconsistent.

4. What are two logical entities in the Cassandra data model? 1 / 1 point

- ( ) Keyspaces and primary keys
- ( ) Tables and definitions
- ( ) Keyspaces and schemas
- (●) Tables and keyspaces

Correct
Correct! Tables are logical entities that organize data storage at cluster and node levels (according to a declared schema), and keyspaces are logical entities that each contain one or more tables.

5. Which guideline should you follow for modeling data? 1 / 1 point

- ( ) Build a primary key based on the maximum number of partitions available.

- ○ Choose a partition key that starts answering your query and that also distributes the data unevenly around the cluster.
- ◉ Build a clustering key that helps you reduce the amount of data that needs to be read by ordering your clustering key columns according to your query.
- ○ Build a primary key that, without regard to the number of partitions, reads to answer a specific query.

Correct

Correct! When modeling data, build a clustering key that helps you reduce the amount of data that needs to be read by ordering your clustering key columns according to your query.