

Your grade: 100%

Your latest: **100%** • Your highest: **100%** • To pass you need at least 70%. We keep your highest score.

Next item →

1. Which is a processing unit of Hadoop and an important core component of the Hadoop framework?

1 / 1 point

- ☒ MapReduce
- ☐ Yet Another Resource Negotiator (YARN)
- ☐ Hadoop Distributed File System (HDFS)
- ☐ Hadoop Common

✓ **Correct**

Correct! MapReduce is a processing unit of Hadoop. It processes data by splitting large amounts into smaller units and processing them simultaneously.

2. Which of the following components are included in Hadoop? Select all that apply.

1 / 1 point

☒ MapReduce

✓ **Correct**

Correct! MapReduce is used to make big data manageable by processing them in clusters.

☐ Apache Cassandra

☒ Hadoop Distributed File System (HDFS)

✓ **Correct**

Correct! HDFS is a core component that stores the data collected from the ingestion and distributes the data across multiple nodes.

☒ Yet Another Resource Negotiator (YARN)

✓ **Correct**

Correct! YARN is the resource manager across clusters.

3. What is the default block size in Hadoop?

1 / 1 point

- ☐ 132 megabytes
- ☐ 126 megabytes
- ☒ 128 megabytes
- ☐ 200 megabytes

✓ **Correct**

Correct! The default block size in Hadoop is 128 megabytes. This block size determines how data is divided and stored across the Hadoop cluster.

4. Which statement is true regarding the comparison between traditional RDBMS and Apache Hive?

1 / 1 point

- ☐ Traditional RDBMS always have built-in support for data partitioning, whereas Hive does not support partitioning.
- ☐ Traditional RDBMS is based on the write once, read many methodologies. Hive allows for as many read operations and write operations as a user needs.
- ☒ Traditional RDBMS can handle up to terabytes of data. Hive is designed to handle petabytes of data.
- ☐ Traditional RDBMS is used to maintain a data warehouse. Hive is used to maintain a database and uses the structured query language known as SQL.

✓ **Correct**

Correct! Hive can handle petabytes of data in the place of terabytes handled by traditional RDBMS.

5. Which component of HBase is a centralized service for maintaining configuration information to maintain healthy links between nodes?

1 / 1 point

- ☒ ZooKeeper
- ☐ Region Servers
- ☐ HMaster
- ☐ Region

✓ **Correct**

Correct! ZooKeeper is a centralized service for maintaining configuration information to maintain healthy links between nodes.

6. Which of the following statements is true with reference to Hive?

1 / 1 point

- ☐ ODBC clients allow Java applications to connect to Hive.
- ☐ JDBC clients allow application based on ODBC to connect to Hive.
- ☐ ODBC client allows applications based on JDBC clients to connect to Hive.
- ☒ The JDBC client allows Java-based applications to connect to Hive

✓ **Correct**

Correct! There are two Hive clients. JDBC clients allows Java based applications to connect to Hive.

7. Which of the following is a feature of Hadoop Distributed File System (HDFS)?

1 / 1 point

- ☐ If one machine crashes, the data needs to be rebuilt again
- ☒ One cluster can be scaled into hundreds of nodes
- ☐ Can store up to megabytes of data
- ☐ Needs permissions to move across multiple platforms

✓ **Correct**

Correct! HDFS is scalable, that is, one cluster can be scaled into hundreds of nodes.

8. What is Yet Another Resource Navigator (YARN)?

1 / 1 point

- ☐ Storage layer in Hadoop
- ☐ Data processing framework
- ☒ Resource Manager
- ☐ Data migration tool

✓ **Correct**

Correct! Yet Another Resource Negotiator (YARN) is the resource management layer in Hadoop, responsible for managing and allocating resources to applications running on a Hadoop cluster.